



Prof. Dr. Stefan Funken
Prof. Dr. Karsten Urban
Lewin Ernst

High Performance Computing II
Institute of Numerical Mathematics
Sheet 1, 34 Points

Exercises in High Performance Computing II *

Introduction

1. We start leisurely by considering the following tasks.

(a) Don't use C-code for the answer.

```
1 // 1. Get the value of a and b as well as of *c, *d and **e.
2 // 2. What is the value of c, d and e?
3 int a = 2, b = 5, *c = &a, *d = &b;
4 int* e[1] = {&a};
5 a = *c * *d;
6 d = c;
7 *d *= *c * *d;
8 a *= **e = 10;
9
```

- (b) Have a look at the material and make yourself familiar with the **structs**, **functions**, **data types** and **utilities**.
- (c) Read [this brief introduction](#) to **makefiles**.
- (d) Run the **makefile** and explain the output.

(4+Karma-Point+Karma-Point+2 = 6 Punkte + 2 Karma-Points)

First impressions of storage formats, algorithms

2. In this exercise we want to get a first impression why storage formats are crucial. We concern ourselves with the storage formats from the lecture. As an example, we introduce a first simple sparse matrix storage format, namely the C00 (coordinate) format, which is also known as the *triplet format*. If a matrix M is given in the C00 format, then

- $M \rightarrow m$; denotes the number of rows.
- $M \rightarrow n$; denotes the number of columns.
- $M \rightarrow nz$; is the number of entries.
- $M \rightarrow nzmax$; is the number of numerical values which is greater or equal to nz . The reason is that we allocate a puffer to store at least nz values.
- $Mx = M \rightarrow x$; is an array, which contains the numerical values of M and has the length $nzmax$.
- $M \rightarrow p$; is an array, which contains the column indices for Mx and has the length $nzmax$.
- $M \rightarrow ind$; is an array, which contains the row indices for Mx and has the length $nzmax$.

The C00 format is not ordered and permits duplicate entries. If the matrix gets larger during runtime, the new entries are appended to the data array. For instance, the data

$$\begin{aligned} m &= 3, \quad n = 3, \quad nzmax = 8, \\ Mx &= [1.0 \ 9.0 \ 2.0 \ 5.4 \ 5.5 \ * \ * \ *], \\ p &= [0 \ 0 \ 1 \ 2 \ 0 \ * \ * \ *], \quad ind = [0 \ 1 \ 2 \ 1 \ 2 \ * \ * \ *], \quad nz = 5. \end{aligned}$$

* Generally: Results are to be explained and code is to be commented.

gives the matrix M

$$M = \begin{pmatrix} 1 & 0 & 0 \\ 9 & 0 & 5.4 \\ 5.5 & 2 & 0 \end{pmatrix}. \quad (1)$$

In C, we combined the matrix formats `COO`, `CRS` and `CCS` in one type `cs`.

Explain the following sparse matrix storage formats:

- (a) `CRS` format (*Compressed Row Storage*), `CCS` format (*Compressed Column Storage*),
- (b) `SED` format (*Sparse Extracted Diagonal*), also called Modified Column Storage,
- (c) `CDS` format (*Compressed Diagonal Storage*),
- (d) `JDS` format (*Jagged Diagonal Storage*),
- (e) `SKY` format (*Skyline Storage*).

What are the advantages of each format?

(2+4+2+2+2 = 12 Punkte)

Parallel matrix vector product

3. In this exercise we will implement the parallel matrix vector product and in the course of this we revise some mpi functions, e.g. `MPI_Bcast` or `MPI_Scatterv`. In order to do so, let $A \in \mathbb{R}^{m \times n}$ be a dense (row major) matrix, $x \in \mathbb{R}^n$ and $y \in \mathbb{R}^m$ vectors. We want to compute

$$y \leftarrow \beta y + \alpha Ax, \quad \alpha, \beta \in \mathbb{R}.$$

To develop our parallel algorithm, we have to think about data decomposition. There are a variety of ways to partition the matrix and vector elements. Each data decomposition results in a different parallel algorithm. We will use a straightforward way: the rowwise block striping. With this strategy, each of the p processes is responsible for a contiguous group of either $\lfloor m/p \rfloor$ or $\lceil m/p \rceil$ rows of the matrix A . The vectors y and x are replicated, meaning all the vector elements are copied on all of the nodes. Each process then calculates inner products of their respective rows of A and x , multiplies the results with α and add it to the scaled entries of y . After the inner product computation each process has some elements of the result y . However, the vectors are supposed to be replicated so that each process has to communicate its results to all other processes (Hint: there is a mpi function for this task: [openmpi doc](#)).

Parallel matrix vector multiplication:

- (a) Why is it acceptable to store the vectors x and y in their entirety on each node, but not the matrix A ?
- (b) Implement the parallel matrix vector multiplication explained above. Assume that m is dividable by p , so that `MPI_Scatter` can be used.
- (c) Extend your code such that m do not need to be dividable by p . Therefore `MPI_Scatter` can not be used anymore and a function `slices` is necessary, which calculates for the input `problem_size`, `number_of_processes` and `process_rank` the `offset` and `size` of the slice (for process `process_rank`).
- (d) Test your functions.

(1+5+6+4 = 16 Punkte)