

1. We start leisurely by considering the following tasks.

(a) Don't use C-code for the answer.

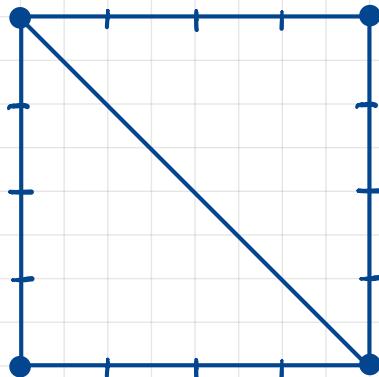
```
1 // 1. Get the value of a and b as well as of *c, *d and **e.  
2 // 2. What is the value of c, d and e?  
3 int a = 2, b = 5, *c = &a, *d = &b; a=2, b=5, c=Zeiger auf(a), d=Zeiger auf(b)  
4 int* e[1] = {&a}; e=Array von Pointern (Zeiger auf(a))  
5 a = *c * *d; a=2*5=10  
6 d = c; d=Zeiger auf(c)  
7 *d *= *c * *d; a=2*2*5=20  
8 a *= **e = 10; a=10*10=100  
9
```

(b) Have a look at the material and make yourself familiar with the **structs**, **functions**, **data types** and **utilities**.

(c) Read [this brief introduction to makefiles](#).

(d) Run the **makefile** and explain the output.

(4+Karma-Point+Karma-Point+2 = 6 Punkte + 2 Karma-Points)



1. Compilieren und linken

2. Führt die Datei `hpc-demo0` aus für zwei Probleme.

In den Problem wird ein großes Netz geladen, welches anschließend verfeinert wird. Wobei die Matrizen im SED (Sparse Extracted Diagonal) Format gespeichert werden.

Zu einer gewählten rechten Seite wird $Ax=b$ näherungsweise mit Gauß-Seidel gelöst und das Residuum berechnet.

First impressions of storage formats, algorithms

2. In this exercise we want to get a first impression why storage formats are crucial. We concern ourselves with the storage formats from the lecture. As an example, we introduce a first simple sparse matrix storage format, namely the COO (coordinate) format, which is also known as the *triplet format*. If a matrix M is given in the COO format, then

- $M \rightarrow m$; denotes the number of rows.
- $M \rightarrow n$; denotes the number of columns.
- $M \rightarrow nz$; is the number of entries.
- $M \rightarrow nzmax$; is the number of numerical values which is greater or equal to nz . The reason is that we allocate a buffer to store at least nz values.
- $Mx = M \rightarrow x$; is an array, which contains the numerical values of M and has the length $nzmax$.
- $M \rightarrow p$; is an array, which contains the column indices for Mx and has the length $nzmax$.
- $M \rightarrow ind$; is an array, which contains the row indices for Mx and has the length $nzmax$.

The COO format is not ordered and permits duplicate entries. If the matrix gets larger during runtime, the new entries are appended to the data array. For instance, the data

$$m = 3, \quad n = 3, \quad nzmax = 8, \\ Mx = [1.0 \ 9.0 \ 2.0 \ 5.4 \ 5.5 \ * \ * \ *], \\ p = [0 \ 0 \ 1 \ 2 \ 0 \ * \ * \ *], \quad ind = [0 \ 1 \ 2 \ 1 \ 2 \ * \ * \ *], \quad nz = 5.$$

* Generally: Results are to be explained and code is to be commented.

gives the matrix M

$$M = \begin{pmatrix} 1 & 0 & 0 \\ 9 & 0 & 5.4 \\ 5.5 & 2 & 0 \end{pmatrix}. \quad (1)$$

In C, we combined the matrix formats COO, CRS and CCS in one type `cs`.

Explain the following sparse matrix storage formats:

- CRS format (*Compressed Row Storage*), CCS format (*Compressed Column Storage*),
- SED format (*Sparse Extracted Diagonal*), also called Modified Column Storage,
- CDS format (*Compressed Diagonal Storage*),
- JDS format (*Jagged Diagonal Storage*),
- SKY format (*Skyline Storage*).

What are the advantages of each format?

(2+4+2+2+2 = 12 Punkte)

a) CRS (Compressed Row Storage) / CCS (Compressed Column Storage)

Bsp.: CRS

$$\begin{pmatrix} 1 & 0 & 0 \\ 9 & 0 & 5.4 \\ 5.5 & 2 & 0 \end{pmatrix} \rightarrow$$

1	9	5,4	5,5	2
0	0	2	0	1
0 1 3				

Speichere 3 Arrays

1. Array - speichert Zeilen - weise alle nicht Null
Einträge

2. Array - speichert die Spalten-Indices zu den Einträgen aus dem 1. Array

3. Array - speichert die Row-Offsets

Pro:

- Umwandlung kann in linearer Komplexität durchgeführt werden

G) SED (Sparse Extracted Diagonal)

Bsp.:
$$\begin{pmatrix} 11 & 12 & 0 & 14 \\ 21 & 22 & 23 & 24 \\ 31 & 0 & 33 & 0 \\ 41 & 0 & 0 & 44 \end{pmatrix} \rightarrow$$

values	11	22	33	44	*	21	31	41	12	23	14	24	□
ptr-ind	0	1	2	3	4	5	6	7	8	9	10	11	12

Speichere zwei Vektoren

1. Vektor - real/double Vektor für Matrix Einträge

2. Vektor - index Vektor

Pro:

- Umwandlung kann in linearer Komplexität durchgeführt werden.

- Diagonal Einträge sind direkt zugänglich (Jacobi-Methode)

c) CDS (Compressed Diagonal Storage)

Bsp.:

$$\begin{pmatrix} 11 & 12 & 0 & 0 & 0 \\ 21 & 22 & 23 & 0 & 0 \\ 0 & 32 & 33 & 34 & 0 \\ 0 & 0 & 43 & 44 & 45 \\ 0 & 0 & 0 & 54 & 55 \end{pmatrix} \longrightarrow \begin{array}{|c|c|c|} \hline * & 11 & 12 \\ \hline 21 & 22 & 23 \\ \hline 32 & 33 & 34 \\ \hline 43 & 44 & 45 \\ \hline 54 & 55 & * \\ \hline \end{array} \quad \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline \end{array}$$

Speicherung der Diagonalen und Nebendiagonale mit Indizes:

-1 für erste Subdiagonale, -2 für zweite usw.

0 für Diagonale

1, 2 für obere Diagonale

Pro:

- Dieses Format ist für nicht symmetrische, diagonal konstruierte Matrizen vor gesehen.

z.B. für strukturierte Matrizen von der finite Elemente Diskretisierung eines Tensor-Produkt-Gitters

d) JDS (Jagged Diagonal Storage)

Bsp.:

$$\begin{pmatrix} 1 & 2 & 3 & 4 & 5 \\ 1 & 6 & 0 & 9 & 5 & 0 \\ 2 & 0 & 0 & 0 & 0 & 3 \\ 3 & 0 & 2 & 0 & 0 & 7 \\ 4 & 0 & 0 & 3 & 2 & 1 \\ 5 & 0 & 0 & 0 & 4 & 8 \end{pmatrix}$$

→

shift left	6	9	5
3	3	0	0
2	2	7	0
3	3	2	1
4	4	8	0

Column pos.

1	3	4
5	0	0
2	5	0
3	4	5
4	8	0

6	3	5
3	2	1
2	7	0
4	8	0
3	0	0

1	3	4
3	2	5
2	4	8
4	5	0
3	0	0

→ val 6 3 2 4 3 9 2 7 8 5 1

col_ind 1 3 2 4 5 3 4 5 8 4 5

index 1 6 10 12

permutation 1 4 3 5 2

Erst werden alle Einträge nach links geschoben und Spalten Position Notiert.

Dann wird der Länge nach Permutiert.

Zum Schluss wird alles in 4 Arrays gespeichert.

Pro:

- Optimiert für sparse Matrix Vektor Multiplikation
- Optimiert für Vektor-Prozesse und allgemeiner SIMD-Maschinen

e) SKY (Skyline Storage)

Bsp.:

$$\begin{pmatrix} 11 & 4 & 12 & 0 & 0 \\ 4 & 22 & 0 & 0 & 41 \\ 12 & 0 & 33 & 89 & 0 \\ 0 & 0 & 89 & 44 & 0 \\ 0 & 41 & 0 & 0 & 55 \end{pmatrix}$$

Diagonal

11 22 33 44 55

Off-Diagonal

4 12 0 89 41 0 0 □

First Column Index

0 1 3 4 7

Speichere 3 Arrays

1.Array - speichert die Diagonalen Einträge ab

2.Array - speichert Zeilen - weise alle Einträge ab

dem ersten nicht-Null - Eintrag bis zur Diagonale ab

3.Array - speichert die Spalten - Offset

Pro:

- Bei Cholesky-Zerlegung vorteilhaft (Form wird beibehalten)
- effizienter Zugriff auf Diagonalen Einträge

Parallel matrix vector product

3. In this exercise we will implement the parallel matrix vector product and in the course of this we revise some mpi functions, e.g. `MPI_Bcast` or `MPI_Scatterv`. In order to do so, let $A \in \mathbb{R}^{m \times n}$ be a dense (row major) matrix, $x \in \mathbb{R}^n$ and $y \in \mathbb{R}^m$ vectors. We want to compute

$$y \leftarrow \beta y + \alpha Ax, \quad \alpha, \beta \in \mathbb{R}.$$

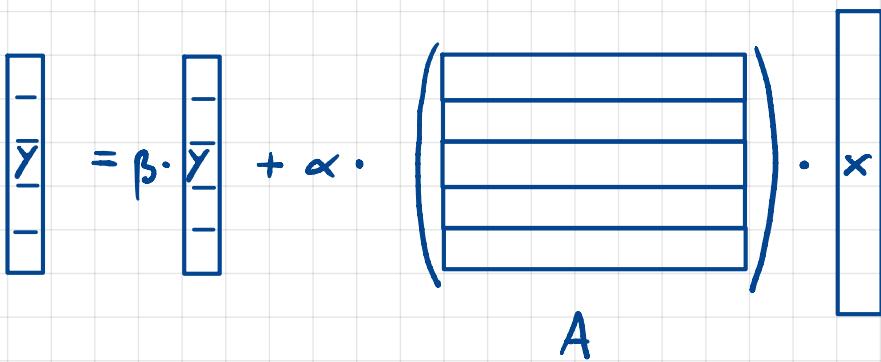
To develop our parallel algorithm, we have to think about data decomposition. There are a variety of ways to partition the matrix and vector elements. Each data decomposition results in a different parallel algorithm. We will use a straightforward way: the rowwise block striping. With this strategy, each of the p processes is responsible for a contiguous group of either $\lfloor m/p \rfloor$ or $\lceil m/p \rceil$ rows of the matrix A . The vectors y and x are replicated, meaning all the vector elements are copied on all of the nodes. Each process then calculates inner products of their respective rows of A and x , multiplies the results with α and add it to the scaled entries of y . After the inner product computation each process has some elements of the result y . However, the vectors are supposed to be replicated so that each process has to communicate its results to all other processes (Hint: there is a mpi function for this task: [openmpi doc](#)).

Parallel matrix vector multiplication:

- (a) Why is it acceptable to store the vectors x and y in their entirety on each node, but not the matrix A ?
- (b) Implement the parallel matrix vector multiplication explained above. Assume that m is dividable by p , so that `MPI_Scatter` can be used.
- (c) Extend your code such that m do not need to be dividable by p . Therefore `MPI_Scatter` can not be used anymore and a function `slices` is necessary, which calculates for the input `problem_size`, `number_of_processes` and `process_rank` the `offset` and `size` of the slice (for process `process_rank`).
- (d) Test your functions.

(1+5+6+4 = 16 Punkte)

a)



x muss im ganzen in jedem Knoten abgespeichert werden, da ganz x für die Berechnung in jedem Knoten gebraucht wird.

Es ist ok x und y in Gänze in jedem Knoten abzuspeichern, da sie bei Skalierung des Problems weniger anwachsen als A .